

Workbook

Zoe Wall
40182161@napier.ac.uk
Edinburgh Napier University - Fundamentals of Parallel Systems (SET09109)

Exercise 2

2.1 Multiplier Process

Code

Listing 1: "Multiplier.groovy"

```
1 void run()
2 {
3     def i = inChannel.read()
4     while (i > 0) {
5         // write i * factor to outChannel
6         outChannel.write(i*factor)
7         // read in the next value of i
8         i = inChannel.read()
9     }
10    outChannel.write(i)
11 }
```

Listing 2: "Consumer.groovy"

```
1 while (i > 0)
2 {
3     //insert a modified println statement
4     println "The output is : ${i}"
5     i = inChannel.read()
6 }
```

Listing 3: "RunMultiplier.groovy"

```
1 def processList = [ new Producer ( outChannel: connect1.out() ),
2
3     //insert here an instance of multiplier with a multiplication factor of 4
4     new Multiplier ( inChannel: connect1.in(),
5                     outChannel: connect2.out(),
6                     factor: 4),
7     new Consumer ( inChannel: connect2.in() )
8 ]
```

Output

```
next: 1
next: The output is : 4
4
next: The output is : 16
10
next: The output is : 40
0
Finished
```

Figure 1: **Exercise 2-1** - Output from Run Multiplier program.

Explanation The *Multiplier.groovy* process (see Listing 1) is inserted into the process list between the producer and consumer. The *Producer.groovy* process outputs an integer, that has been provided, to the Multiplier process which then outputs each integer multiplied by a constant factor which is set in the constructor of the Multiplier instance (see line 6 in Listing 3). The *Consumer.groovy* process prints a meaningful output to the console.

2.2 Integer Sets

Code

Listing 4: "ListToStream.groovy"

```
1 while (inList[0] != -1)
2 {
3     // hint: output list elements as single integers
4     for ( i in 0 ..< inList.size)outChannel.write(inList[i])
5     inList = inChannel.read()
6 }
```

Listing 5: "CreateSetsOfEight.groovy"

```
1 while (v != -1)
2 {
3     for ( i in 0 .. 7 )
4     {
5         // put v into outList and read next input
6         outList[i] = v
7         v = inChannel.read()
8     }
9     println " Eight Object is ${outList}"
10 }
```

Listing 6: "GenerateSetsOfThree.groovy"

```
1 void run()
2 {
3     def threeList = [
4         [1, 2, 3],
5         [4, 5, 6],
6         [7, 8, 9],
7         [10, 11, 12],
8         [13, 14, 15],
9         [16, 17, 18],
10        [19, 20, 21],
11        [22, 23, 24],
12        [-1, -1, -1]] // terminating list
13    for ( i in 0 ..< threeList.size)outChannel.write(threeList[i])
14    //write the terminating List as per exercise definition
15 }
```

Output

```
Eight Object is [1, 2, 3, 4, 5, 6, 7, 8]
Eight Object is [9, 10, 11, 12, 13, 14, 15, 16]
Eight Object is [17, 18, 19, 20, 21, 22, 23, 24]
Finished
```

Figure 2: **Exercise 2-2** - Output from Run Three to Eight program.

Exercise Questions

What change is required to output objects containing six integers?

Within the *CreateSetsOfEight.groovy* process, change the number of iterations of the for loop. See line 3 of listing 5. The new line should read

Listing 7: "CreateSetsOfEight.groovy - Change required to output objects containing six integers"

```
1 //for (i in 0 .. 7) {
2 for (i in 0 .. 5) {
```

which will now create lists containing 6 integers. This can be improved by changing the number to a variable which is set through the constructor of the process.

How could you parameterise this in the system to output objects that contain any number of integers (e.g. 2, 4, 8, 12)?

The process can be improved by creating a parameter for the size of the output list, bearing in mind the iterations start from zero so the required size should be minus one of the input parameter. This can either be a variable that is set within the constructor of the *CreateSetsOfEight.groovy* process or by user input from the console.

What happens if the number of integers required in the output stream is not a factor of the total number of integers in the input stream (e.g. 5 or 7) ?

Some integers from the input stream will not be outputted to the console as they do not make up a full set of numbers. So the set cannot be filled and therefore the set will not display as the process cannot finish.

Exercise 3

3.1 Reversing GIntegrate




Figure 3: **Exercise 3-1** - Network for process reversing the effect of *GIntegrate.groovy*.

GNumbers.groovy outputs a stream of integers starting at zero and incrementing by one each time. *GIntegrate.groovy* increments the stream by an increasing number each time so the difference between the output is increasing. To negate this effect the output from the network should be equal to the output from the initial *GNumbers* process.

Minus Process

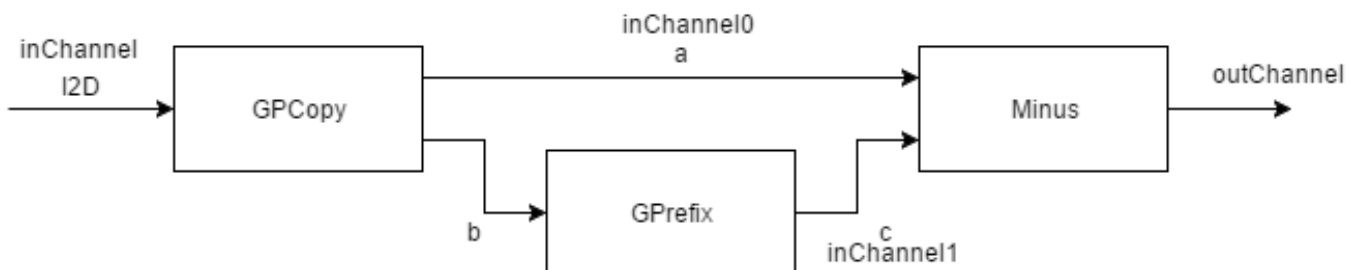


Figure 4: **Exercise 3-1** - Network showing processes for *Differentiate.groovy*.

Explanation The minus process works by copying the output value from *GIntegrate.groovy*, e.g 0, 1, 3, 6. One copy is sent straight to *inChannel0* of the *Minus.groovy* process, and the other is sent to *GPrefix.groovy* which outputs the stream with a leading zero into *inChannel1* - see Figure 4. The *Minus.groovy* process reads both inputs in parallel and minuses the second from the first resulting in an output stream of incrementing numbers (Figure 5).

Code

Listing 8: "Minus.groovy"

```

1 ProcessRead read0 = new ProcessRead ( inChannel0)
2 ProcessRead read1 = new ProcessRead ( inChannel1)
3 def parRead2 = new PAR ( [ read0, read1 ] )
4
5 while (true) {
  
```

```

6    parRead2.run()
7    // output one value subtracted from the other
8    // be certain you know which way round you are doing the subtraction!!
9    outChannel.write(read0.value - read1.value)
10 }

```

Listing 9: "Differentiate.groovy - see Figure 4 for network of this list"

```

1  def differentiateList = [ new GPrefix ( prefixValue: 0,
2                          inChannel: b.in(),
3                          outChannel: c.out() ),
4                          new GPCopy ( inChannel: inChannel,
5                                      outChannel0: a.out(),
6                                      outChannel1: b.out() ),
7                          // insert a constructor for Minus
8                          new Minus (inChannel0: a.in(),
9                                    inChannel1: c.in(),
10                                   outChannel: outChannel)
11 ]

```

Output

```

Differentiated Numbers
0
1
2
3
4
5
6
7
8
9
10

```

Figure 5: **Exercise 3-1** - Output from the Differentiate System using the Minus Process.

Negator Process




Figure 6: **Exercise 3-1** - Network showing processes for *DifferentiateNeg.groovy* (replacement for Differentiate process within Figure 3)

Explanation The negator process works by using the *GPlus.groovy* process which outputs the sum of the two input channels read in parallel. The first input channel is a copy of the original input, whereas the second input channel is a negative version of the input with a leading zero from *GPrefix.groovy*.

Code

Listing 10: "DifferentiateNeg.groovy"

```

1  def differentiateList = [ new GPrefix ( prefixValue: 0,
2                          inChannel: b.in(),
3                          outChannel: c.out() ),
4                          new GPCopy ( inChannel: inChannel,
5                                      outChannel0: a.out(),
6                                      outChannel1: b.out() ),
7                          //insert a constructor for Negator
8                          new Negator ( inChannel: c.in(), outChannel: d.out()),
9                          new GPlus ( inChannel0: a.in(),

```

```

10         inChannel1: d.in(),
11         outChannel: outChannel )
12     ]

```

Output

```

Differentiated Numbers
0
1
2
3
4
5
6
7
8
9
10

```

Figure 7: **Exercise 3-1** - Output from the Differentiate System using the Negator process.

Exercise Questions

Which is the more pleasing solution and why? In this case, even though the Negator solution requires an extra process, it can be argued that this is more pleasing as there is less room for error introduced by the order of the values within the Minus solution. However, a more appropriate solution to negating the effects of GIntegrate would be not to send the values through it in the first place.

3.2 Sequential Copy Process

GSCopy Code

Listing 11: "GSCopy.groovy"

```

1  void run ()
2  {
3      while (true)
4      {
5          def i = inChannel.read()
6          // output the input value in sequence to each output channel
7          outChannel0.write(i)
8          outChannel1.write(i)
9      }
10 }

```

Explanation Within the GSCopy the input value is copied by outputting the value to both output channels in sequence not in parallel.

GSPairsA




Figure 8: **Exercise 3-2** - Process Network diagram of GSPairsA, as the copy is sequential the output is written first to channel **a** then to channel **b**.

Output

Squares

Figure 9: **Exercise 3-2** - Output from the Squares system, using the *GSPairsA.groovy* process.

GSPairsB




Figure 10: **Exercise 3-2** - Process Network diagram of GSPairsA, as the copy is sequential the output is written first to channel **b** then to channel **a**.

Output

Squares
 1
 4
 9
 16
 25
 36
 49
 64
 81
 100

Figure 11: **Exercise 3-2** - Output from the Squares system, using the *GSPairsB.groovy* process.

Questions




Figure 12: **Exercise 3-2** - Process Network diagram of test system, GSPairs is replaced by GSPairsA and GSPairsB to see the effect.

Determine the effect of the change between GSPairsA and GSPairsB, why does this happen?

In running *GSPairsA.groovy* the output, shown in Figure 9, is blank after the header, whereas in running *GSPairsB.groovy* the output, shown in Figure 11, displays a heading and then a stream of square numbers. The reason for this change is due to the fact *GSCopy.groovy* outputs sequentially and even though the channels are the

same, the order in which they are written to is different. See Figures 8 & 10. In the instance using *GSPairsA.groovy* the system deadlocks. *GPlus.groovy* reads both inputs in parallel so cannot process the values until both channels have an input. *GTail.groovy* removes the first input it is given and then outputs the remaining numbers. However *GTail.groovy* cannot receive another input until channel a (the channel between *GSCopy.groovy* and *GPlus.groovy*) is free. The instance using *GSPairsB* does not deadlock as channel b - the channel between *GSCopy.groovy* and *GTail.groovy* - is written to first, so will be able to receive and send the next copied value on the next iteration of the sequential copy.

3.3 Parallel Print

Questions

Why was it considered easier to build GParPrint as a new process rather than using multiple instances of GPrint to output the table of results?

It is easier to build a parallel printing process rather than using multiple instances of *GPrint* as it allows for printing results from multiple processes at once. It also allows for However as a parallel print needs to wait for every input process to return a value to print, it means that it can only output as fast as the slowest process. Therefore it would only be easier to use if the processes to print output at a regular rate.

Exercise 4

4.1 ResetPrefix

Listing 12: "Line 25 of ResetPrefix.groovy"

```
1 inChannel.read()
```

Output

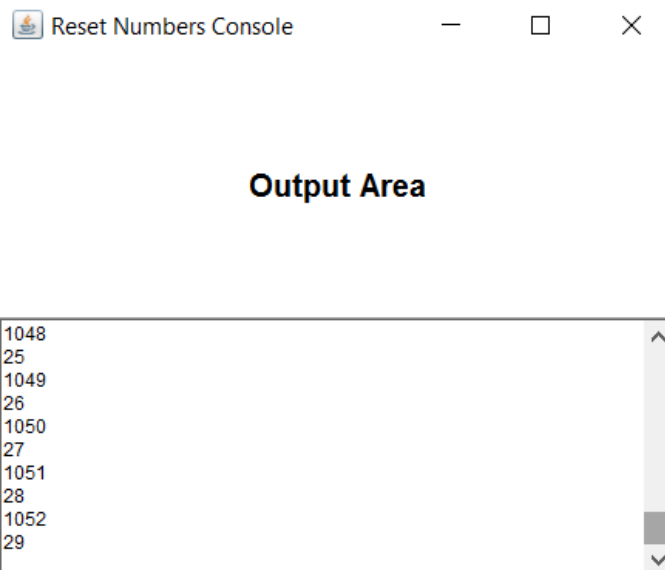


Figure 13: **Exercise 4-1** - Output from the ResetNumbers program when line 25 of *ResetPrefix.groovy* is removed.

Questions

What happens if line 25 of ResetPrefix Listing 4-1 is commented out? Why?

When the read line is removed the output alternates between the incrementing reset value and the original numbers as seen in figure 13. This is because the original value is not removed from the system if the channel is not read from.

Explore what happens if you try to send several reset values hence, explain what happens and provide a reason for this.

When you try to send several reset values at once deadlock occurs. This is due to the first number not being removed from channel c, shown in figure 14, and therefore the processes cannot read additional inputs from the resetChannel as there is no way for *GPCopy.groovy* to output in parallel until the channel is cleared.

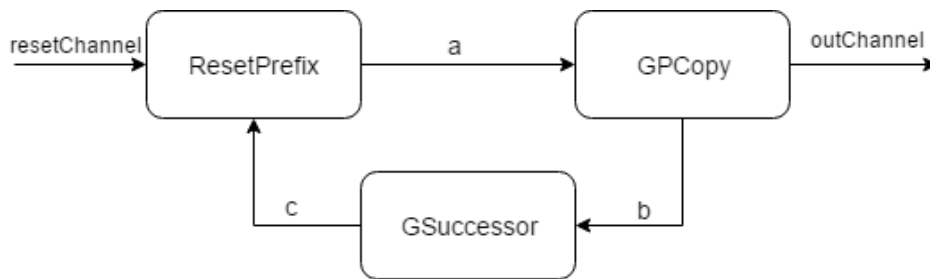


Figure 14: **Exercise 4-1** - Process network diagram for *ResetNumbers.groovy*

4.2 ResetSucessor process

Code

Listing 13: "ResetSucessor.groovy"

```

1 while (true)
2 {
3     // deal with inputs from resetChannel and inChannel
4     // use a priSelect
5     def index = alt.priSelect();
6
7     if (index == 0)    // reset Channel input
8     {
9         def resetVal = resetChannel.read()
10        // inChannel.read()
11        outChannel.write(resetVal)
12    }
13    else                // outChannel input
14    {
15        outChannel.write(inChannel.read() + 1)
16    }
17 }

```

Listing 14: "ResetNumbers.groovy"

```

1 def testList = [
2     new GPrefix ( prefixValue: initialValue,
3         outChannel: a.out(),
4         inChannel: c.in() ),
5     new GPCopy ( inChannel: a.in(),
6         outChannel0: outChannel,
7         outChannel1: b.out() ),
8     // requires a constructor for ResetSucessor
9     new ResetSucessor (inChannel: b.in(),
10         outChannel: c.out(),
11         resetChannel: resetChannel )
12 ]

```

Question

Does it overcome the problem identified in Exercise 1? If not, why not?

No, the reformulation does not overcome the problem as it is still possible for the user to input multiple values, deadlocking the system. There is no way to reformulate this example without reading from the channel as the integers are never removed from the system.

Exercise 5

5.1 Varying delay for RunQueue

Questions *By varying the delay times demonstrate that the system works in the manner expected. What do you conclude from these experiments?*

Varying the delay times makes no difference to the process output, it works as expected regardless of changing the times. This happens because it uses preconditions to

5.2 Preconditions

Code

Listing 15: "Scale.groovy"

```
1 while (true)
2 {
3     switch ( scaleAlt.priSelect(preCon) )
4     {
5         case SUSPEND :
6             // deal with suspend input
7             suspend.read()
8             factor.write(scaling)
9             suspended = true
10            println "Suspended"
11            preCon[SUSPEND] = false
12            preCon[INJECT] = true
13            break
14        case INJECT:
15            // deal with inject input
16            scaling = injector.read()
17            println "Injected scaling is $scaling"
18            suspended = false
19            timeout = timer.read() + DOUBLE_INTERVAL
20            timer.setAlarm(timeout)
21            preCon[SUSPEND] = true
22            preCon[INJECT] = false
23            suspended = false
24            break
25        case TIMER:
26            // deal with Timer input
27            timeout = timer.read() + DOUBLE_INTERVAL
28            timer.setAlarm ( timeout )
29            scaling = scaling * 2
30            println "Normal Timer: new scaling is ${scaling}"
31            break
32        case INPUT:
33            // deal with Input channel
34            def inValue = inChannel.read()
35            def result = new ScaledData()
36            result.original = inValue
37            result.scaled = inValue * scaling
38            outChannel.write( result )
39            break
40    } //end-switch
41 } //end-while
```

Questions

Which is the more elegant formulation? Why?

The latter, using preconditions, is the more elegant solution because it avoids nested loops. Using only one switch statement and alternative makes the code more readable and possibly could increase performance due to a decrease in conditional loops.

Exercise 6

6.1 Test Case for Three-To-Eight

Code

Listing 16: "RunThreeToEightTest.groovy"

```
1 class RunThreeToEightTest extends GroovyTestCase
2 {
3     void testThreeToEight()
4     {
5         One2OneChannel genToStream = Channel.one2one()
6         One2OneChannel streamToEight = Channel.one2one()
```

```

7
8  def gen = new GenerateSetsOfThree ( outChannel: genToStream.out())
9  def list = new ListToStream ( inChannel: genToStream.in(), outChannel: streamToEight.out())
10 def eight = new CreateSetsOfEight ( inChannel: streamToEight.in())
11
12  def testRunList = [gen, list, eight]
13  new PAR(testRunList).run()
14
15  // test output is correct from eight
16  def expectedList = list.inTest
17  def actualList = eight.outTest
18
19  println "exp ${expectedList} + act ${actualList}"
20
21  assertTrue(expectedList == actualList)
22
23  assertTrue(eight.outList.size() == 8)
24 }
25 }

```

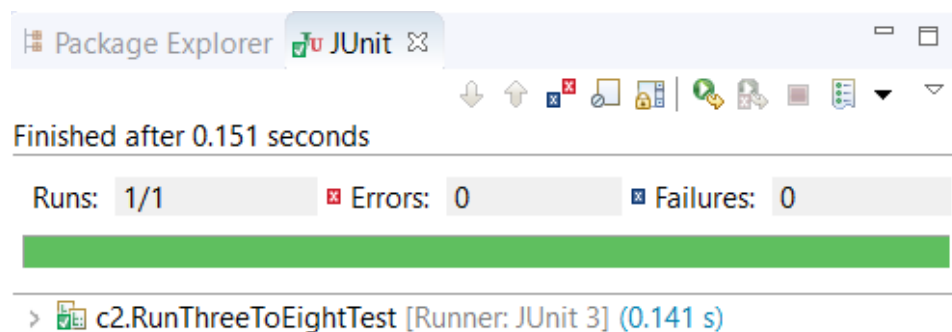


Figure 15: **Exercise 6-1** - Screen capture of the results of the JUnit test performed.

Output

Explanation The performed JUnit test shown above adds each integer from the *ListToStream.groovy* process in turn and checks it against the output from *CreateSetsOfEight.groovy* to check if the values are correctly being processed through the system. It also checks that the set returned from the *CreateSetsOfEight.groovy* process correctly contains 8 integers.

Exercise 7

7.1 Deadlock

Determine the precise nature of the deadlock in the Client Server system.

By adding a parameter to identify each server you can see from the output of the program, figure 16, that both servers request values from each other at the same time which is the cause of the deadlock in the system. Neither process can send a response as they first are expecting a response for their own request and therefore cannot continue.

Output

```

Client number 1 requests 14
Server 1 recieves request from other server.
server 0 gets request from server
Server 1 recieves request for 14 from CLIENT
Client number 0 requests 3
Client number 1 requests 15
send 0 as a client
Server 0 recieves request for 3 from CLIENT
Server 1 recieves request for 15 from CLIENT
Client number 0 requests 14
Client number 1 requests 6
Server 0 recieves request for 14 from CLIENT
Server 1 recieves request for 6 from CLIENT
Server 1 requests other server for 6
Server 0 requests other server for 14

```

Figure 16: **Exercise 7-1** - Output from the Server Client process, deadlock occurs as each server requests a response from the other.

Exercise 8

8.1 Test Case for Client Server System

Code

Listing 17: "Client.groovy"

```

1 void run()
2 {
3     def iterations = selectList.size
4     println "Client $clientNumber has $iterations values in $selectList"
5
6     for ( i in 0 ..< iterations) {
7         def key = selectList[i]
8         println "Client number $clientNumber requests $key"
9         requestChannel.write(key)
10        def v = receiveChannel.read()
11
12        // add response from server to actual response list
13        actualList << v
14    }
15
16    println "Client $clientNumber has finished"
17
18    // multiply each value by ten and add in order from selectList = expected value from server
19    for(i in 0 ..< iterations)expectedList << selectList[i]*10
20
21    // check if actual equals expected
22    if (actualList.equals(expectedList))
23        println "test passed"
24    else
25        println "test failed"
26 }

```

```

Client 1 has finished
Client 0 has finished
test passed
test passed

```

Figure 17: **Exercise 8-1** - Output from the Server Client test.

Output

Explanation To ensure that the values returned from the Server arrive in the order expected according to their selectList property, a list was created to store the actual output from the server. The selectList was then iterated through multiplying each value by ten and storing the result into a separate list. These lists were compared to test the system.

Exercise 9

9.1 Missed Test

Explanation Using the suggestion (Section 9.4.4) made earlier in the chapter, construct an additional process for the event handling system that ensures that the number of missed events is correct. The additional process should be added to the network of processes. You may need to modify the EventData class (Section 9.2.4) to facilitate this.

In constructing a test case for the event handling system, a new *EventValidation.groovy* process was implemented, see Listing 18. The process was inserted into the system between the *UniformDistributedDelay.groovy* and the *GPrint.groovy* process.

To test if the number of missed events is correct, the previous EventData object to pass through the system is stored and updated each loop. EventData instances store the number of missed events that are printed to the console. The algorithm checks whether the missed event stored by the current EventData object is correct by checking against a calculation.

Code

Listing 18: "EventValidation.groovy"

```
1 class EventValidation implements CSProcess
2 {
3     def ChannelInput inChannel
4     def ChannelOutput outChannel
5     def EventData previous
6
7     void run()
8     {
9         while (true)
10        {
11            // read in current event data
12            def currentEvent = (EventData)inChannel.read()
13
14            // previous is null for first event
15            if (previous != null)
16            {
17
18                def actualMissed = currentEvent.data - previous.data - 1
19
20                // check if current missed event value is the same as the calculated missed value
21                if (currentEvent.missed != actualMissed)
22                {
23                    println "Error: Missed count inaccurate.\n"
24                    +"Event data: $currentEvent.data missed count = $currentEvent.missed, should be $actualMissed"
25                }
26            }
27
28            // send along to GPrint and set new previous event
29            outChannel.write(currentEvent)
30            previous = currentEvent
31        } // end while
32    }
33 }
34
35
36 }
```

9.2 MultiStream

Exercise 9 2 (5.marks) The accompanying exercise package contains a version of the event handling system, Run-MultiStream, which allows the creation of 1 to 9 event streams. By modifying the times associated with each event generation stream and also of the processing system explore the performance of the system. What do you conclude?

Through experimentation it can be concluded that there is a direct correlation between the generation/processing delay and the number of missed events. The lower the delay on the event processor, the

fewer events will be missed.

In direct comparison, the lower the delay on the event generation, the more events will be missed. This is because events are constantly written into the EventProcessor, even if it is not ready to receive data. If events are being generated in quick succession it means the EventProcessor cannot keep up with demand and events will be missed.

A solution to this would be to implement a buffer or queue process, similar to the one featured in Exercise 5-1

Table 1: A sample of the correlation between Generation/Processing times and Missed Events

Average Generation Time	Average Processing Time	Missed Events
-------------------------	-------------------------	---------------

175	25	0
-----	----	---

175	175	190
-----	-----	-----

25	175	276
----	-----	-----

Note: All samples were generated using three Event Sources.

By peaking the code definitions and viewing the code sources as well as sample testing, it was a simple task to determine the workings of the three Multiplexers and their methods of channel selection.

FairMultiplex Uses a round-robin selection to try and ensure each channel is given equal opportunity to have its event read. Once a channel has been read it is given the lowest priority, making it unlikely to be read again before another channel.

PriMultiplex Prioritizes a single channel to its fruition, then selects a new channel. This means that the first selected channel has very few missed events but the later selections will each have exponentially higher number of misses.

Multiplexer Actually utilizes the same selection method as the FairMultiplex but under a different name. This arguably, makes it the least useful, as it is simply the FairMultiplex, but with less detailed documentation.

14

SET09109: Workbook 40056761

Summay With the following points in mind, it is clear that the FairMultiplexer is the best choice the majority of the time. In cases where you only care about a single processes output, a PriMultiplexer may be a better choice but in such a case why even use multiple processes. There appears to be no need to ever use the third Multiplex as it simply the FairMultiplex under a different name.

9.3 multiplexer

Exercise 9 3 (6 marks) The process EventProcessing has three versions of multiplexer defined within it, two of which are commented out. By choosing each of the options in turn, comment upon the effect that each multiplexer variation has on overall system performance.

Exercise 11

11.1 control

Exercise 11 1 (8 marks) The Control process in the Scaling system currently updates the scaling factor according to an automatic system. Replace this with a user interface that issues the suspend communication, obtains the current scaling factor and then asks the user for the new scaling factor that is then injected into the Scaler. The original and scaled values should also be output to the user interface. Total Marks (60)